

CropDoc

AI-powered plant care

Chandrabhushan Upadhyay



Introduction

The agricultural sector, from large-scale farming operations to small home gardens, faces a constant threat from plant diseases. The effective management of crop health is a critical factor in ensuring food security, potentially increasing yields, and reducing the unnecessary use of pesticides. For both commercial farmers and individual gardeners, the ability to quickly and accurately respond to the first signs of illness in a plant can be the difference between a successful harvest and significant loss.

While agricultural experts provide invaluable knowledge, their services can be costly and are not always readily available to those who need them most. This accessibility issue presents the central challenge in modern plant health management. The absence of an immediate, affordable, and dependable method for diagnosis creates a critical gap for growers. This gap between the onset of a disease and its accurate identification is the origin point for many preventable losses, leading directly into the core problem that this project aims to solve.

Problem Statement

The agricultural sector, from large-scale commercial farming operations to individual home gardens, faces a constant and significant threat from plant diseases. The health and vitality of crops are fundamental to both economic stability and food security, yet the process of accurately identifying and managing plant illnesses remains a persistent and critical challenge for growers at every level.

The primary barrier to effective plant health management is the lack of an accessible, real-time, and reliable method for diagnosis. Traditionally, identifying a disease requires a high level of specialized knowledge, often necessitating consultation with agricultural experts whose services can be costly and are not always readily available. Existing solutions are

frequently manual, time-consuming, and expensive, creating a crucial delay between the initial detection of symptoms and the application of an effective treatment. This gap leaves crops vulnerable and puts growers at a significant disadvantage.

This delay and the high potential for misidentification lead to severe consequences. Applying an incorrect treatment is not only ineffective but directly contributes to crop loss, reduced yields, and the wasteful expenditure of valuable resources like water and fertilizer. Furthermore, in an attempt to combat an unknown ailment, growers may resort to the broad or unnecessary use of chemical pesticides, which increases operational costs and can contribute to environmental harm.

Strategic Analysis

Feasibility

The CropDoc service demonstrates high **short-term feasibility**. The project already exists as a functional proof-of-concept with a defined technology stack and a validated CNN model that exceeds its performance targets. The final report outlines a clear, well-defined next phase with a modest budget of approximately \$10,000 and a 6-month schedule. This plan includes expanding the dataset and developing native mobile apps, showing that the core product and its subsequent commercial version can be developed and launched well within the 2-3 year timeframe.

Viability

The product has excellent long-term viability because it addresses a fundamental and perpetual need in the agricultural sector: the "constant threat from plant diseases". As long as food is grown, there will be a need to protect crops. The project's approach of using AI and computer vision aligns with the long-term trend of technological innovation in agriculture. This ensures the service will remain relevant and can evolve with technological advancements, making it a sustainable solution for the coming decades.

Monetization

The CropDoc platform is designed for **direct monetization** through multiple revenue

streams. The primary business model is built on affiliate marketing and strategic sponsorships, where the platform earns a commission or a referral fee for recommending specific, sponsored cure products. Additionally, the project outlines plans for direct revenue through premium subscription tiers, which would offer advanced analytics to commercial farmers and an ad-free experience to users. A third direct channel includes licensing aggregated, anonymized data to agricultural research institutions.

Code Implementation and model building

The development of the CropDoc prototype involved a carefully selected technology stack and a methodical approach to building and validating the core machine learning model. The implementation was divided into the frontend interface, the backend processing logic, and the deep learning model itself.

Technology Stack

The prototype was built using a modern and efficient set of technologies chosen for their accessibility and performance. The backend was developed using

Python - a powerful language for machine learning, along with the **Flask** web framework to handle image uploads and API requests. The core of the system is a pre-trained

ResNet-50 model, which was trained and deployed using the **TensorFlow** library.

Flask API

The Flask API serves as the central backend component of the CropDoc application, acting as the crucial bridge between the user-facing frontend and the machine learning model. Its primary responsibilities within the workflow include:

- **Receiving Image Uploads:** It handles the HTTP request from the user's browser, securely receiving the uploaded plant leaf image.
- **Orchestrating the Prediction:** The API preprocesses the image to meet the model's input requirements and then passes it to the pre-trained ResNet-50 model to get a disease prediction.

- **Retrieving Treatment Data:** Using the model's prediction, it queries the Firestore database to fetch the corresponding disease name, description, and recommended cure.
- **Sending the Final Response:** Finally, it packages all the relevant information into a structured format and sends it back to the frontend to be displayed to the user in under 5 seconds.

```

1 # Import necessary libraries
2 from flask import Flask, render_template, request
3
4 import numpy as np
5 import os
6
7 from tensorflow.keras.preprocessing.image import load_img
8 from tensorflow.keras.preprocessing.image import img_to_array
9 from tensorflow.keras.models import load_model
10
11 filepath = r"C:/Users/Chandrabhushan/.vscode/Project from scratch/ Tomato Leaf Disease Detection/model.h5"
12
13 model = load_model(filepath)
14 print(model)
15
16 print("Model Loaded Successfully")
17
18 def predict_tomato_disease(tomato_image):
19     test_image = load_img(tomato_image, target_size=(128, 128)) # Load image
20     print(f"Got image for prediction")
21
22     test_image = img_to_array(test_image) / 255 # convert image to np array and normalize
23     test_image = np.expand_dims(test_image, axis=0) # change dimension 3D to 4D
24
25     result = model.predict(test_image) # predict diseased plant or not
26     print(f"Raw result = ", result)
27
28     pred = np.argmax(result, axis=1)
29     print(pred)
30     if pred == 0:
31         return "Tomato Bacteria Spot Disease", "Tomato-Bacteria Spot.html"
32
33     elif pred == 1:
34         return "Tomato Early Blight Disease", "Tomato-Early_Blight.html"
35
36     elif pred == 2:
37         return "Tomato Healthy and Fresh", "Tomato-Healthy.html"
38
39     elif pred == 3:
40         return "Tomato Late Blight Disease", "Tomato - Late_blight.html"
41
42     elif pred == 4:
43         return "Tomato Leaf Mold Disease", "Tomato - Leaf_Mold.html"
44
45     elif pred == 5:
46         return "Tomato Septoria Leaf Spot Disease", "Tomato - Septoria_leaf_spot.html"
47
48     elif pred == 6:
49         return "Tomato Target Spot Disease", "Tomato - Target_Spot.html"
50
51     elif pred == 7:
52         return "Tomato Yellow Leaf Curl Virus Disease", "Tomato - Tomato_yellow_Leaf_Curl_Virus.html"
53
54     elif pred == 8:
55         return "Tomato Mosaic Virus Disease", "Tomato - Tomato_mosaic_virus.html"
56

```

Figure 1

Model Architecture and Training

The heart of **CropDoc** is a **Convolutional Neural Network (CNN)**, which is the industry standard for image classification tasks. After considering several architectures, the pre-trained

ResNet-50 model was selected for its balance of high accuracy and cost-effectiveness over building a custom model from scratch.

The model was trained on a dataset of plant images to classify different diseases. The validation process was rigorous, with the primary objective of achieving a prediction accuracy of over 95% on the test dataset. The final trained model successfully surpassed this goal, achieving an **accuracy of 96.2%**, confirming the technical viability of the concept.




Figure 2

Backend Implementation

The backend, built with Python and Flask, serves as the brain of the operation. Its workflow is as follows:

1. A user uploads an image of a plant leaf through the frontend interface. The image is sent to the backend server.
2. The backend script receives the image and preprocesses it by resizing and cleaning it to match the model's required input format.

3. The prepared image is then fed into the ResNet-50 model, which returns a prediction for the disease.
4. This prediction is used as a key to query the Firestore database to retrieve the disease name, a detailed description, and information about the sponsored cure.
5. Finally, all of this information is packaged and sent back to the frontend to be displayed to the user. The entire process was optimized to ensure the user receives a result in under 5 seconds.




Figure 3

Frontend Implementation

The user interface was intentionally designed for simplicity and ease of use, adhering to a "3-click diagnosis" principle. The single-page web application presents the user with a straightforward "Upload Image" button. Once the backend has processed the image, the results are displayed on a clear, easy-to-understand page that shows the diagnosed disease and the recommended treatment, including a link to the sponsored brand's product. This minimalist approach ensures that users of all technical skill levels can access the tool's powerful diagnostic capabilities.




Figure 4




Figure 5

CropDoc Business Model

The business model for CropDoc is designed for sustainability and growth by integrating several direct monetization streams while keeping the core diagnostic tool free for the end-user. This approach ensures wide accessibility for farmers and gardeners, which in turn builds a valuable user base for the platform's commercial strategies.

Primary Revenue: Affiliate Marketing & Strategic Sponsorships

The main revenue stream is generated through a highly targeted affiliate marketing system.

- When the application diagnoses a plant disease, it recommends a specific cure and directs the user to a sponsored product that can solve the problem.
- In this mutually beneficial relationship, the user receives an immediate and actionable solution, while the brand partner gets a qualified lead at the exact moment of need.
- CropDoc earns a commission for each resulting sale or a fixed fee for each referral, creating a sustainable business model without charging the user.

Secondary Revenue: Premium Subscription Tiers

While the core functionality remains free, CropDoc will offer premium features through a subscription model to serve more advanced user needs. These paid tiers are designed for distinct user segments:

- **For Commercial Farmers:** A subscription provides access to advanced data and analytics, including detailed historical data on disease outbreaks, severity trends, and predictive analytics for their specific crops.
- **For General Users:** A paid tier offers customizable content, allowing users to create personalized plant profiles and track health history, along with an ad-free experience that removes sponsored recommendations.

Tertiary Revenue: Data Licensing and Analytics

A third revenue stream involves the licensing of anonymized and aggregated data.

- With user consent, data on plant disease occurrences, geographical distribution, and user behavior can be licensed to agricultural research institutions, pesticide manufacturers, and government agencies.
- This data provides valuable insights into agricultural trends and can inform product development and proactive disease management on a larger scale

Market Identification and Opportunity

The primary market for the launch of **CropDoc** is the burgeoning **Indian Agri-Tech sector**. This is not a single, uniform market but is composed of distinct segments with specific needs. A successful launch requires a nuanced strategy that addresses both.

Target Market Segments

1. **B2B (Business-to-Business): Commercial Farmers and Growers** This segment represents the economic backbone of Indian agriculture. For these users, farming is a business, and their primary concerns are **yield optimization, cost reduction, and operational efficiency**. They face significant financial losses from crop diseases and are increasingly open to technological solutions that provide a clear return on investment. CropDoc's value proposition for them is its ability to provide rapid, accurate diagnostics that can prevent crop loss and reduce the misapplication of expensive pesticides and fertilizers.
2. **B2C (Business-to-Consumer): Home Gardeners and Enthusiasts** This is a rapidly growing segment, especially in urban and semi-urban areas. These users are motivated by hobbies, a desire for organic produce, and the aesthetics of gardening. While they are not focused on large-scale yield, they are willing to spend on products that ensure the health and beauty of their plants. For this segment, CropDoc serves as an accessible expert, providing immediate answers and easy-to-follow, often organic, treatment recommendations.

Key Market Drivers

The timing is ideal for a product like CropDoc due to several converging factors:

- **Digital Transformation:** Government initiatives like 'Digital India' and the increasing affordability of smartphones and data have led to massive digital penetration, even in

rural areas. This makes a mobile-first solution like CropDoc technically feasible and accessible to a vast user base.

- **Focus on Sustainability:** There is a growing awareness of sustainable agricultural practices. CropDoc supports this by enabling targeted treatments, which can reduce the overuse of chemical pesticides, aligning with both environmental goals and farmer priorities.
- **Venture Capital Interest:** The Indian Agri-Tech space is attracting significant investment, indicating strong market confidence and providing potential future funding opportunities for scaling the CropDoc platform.

Market Data and Strategic Implications

Understanding the market statistics is crucial for appreciating the scale of the opportunity.

- **Total Addressable Market (TAM):** With an estimated **150 million farmers** and a growing population of millions of home gardeners, the potential user base is enormous. This vast number represents the Total Addressable Market that CropDoc can tap into.
- **Value of the Problem:** Annual crop losses due to pests and diseases amount to billions of dollars. This figure represents the monetary value of the problem CropDoc is solving, making its value proposition incredibly compelling. A tool that can mitigate even a small fraction of these losses provides immense value.
- **The Digital Enabler:** With over **300 million smartphone users in rural India** alone, the primary barrier to delivering a technological solution—access—has been significantly lowered. This digital readiness is the critical enabler that makes CropDoc's business model viable at scale.

Financial Modelling and Forecasting

To forecast the financial potential of CropDoc, we utilize a revenue projection model. This model is not just a calculation but a strategic tool to test assumptions, understand key growth drivers, and project a realistic financial future.

The Revenue Projection Model

The core of the financial forecast is a direct revenue equation based on the platform's primary monetization strategy: affiliate marketing and sponsorships.

Equation:

Projected Annual Revenue = $U \times F \times CR \times C$

Detailed Variable Analysis

- **U (Number of Active Users):** This is the most critical growth metric. It will be driven by a dedicated user acquisition strategy, moving potential users through a marketing funnel from awareness to active use. Initial growth will be driven by digital marketing, with long-term retention depending on the app's accuracy and user experience.
- **F (Usage Frequency):** This variable represents user engagement. We assume an average of 3 diagnoses per user annually, but this could fluctuate based on seasonal disease outbreaks, the types of crops supported, and the user's level of engagement with their farm or garden.
- **CR (Conversion Rate):** This metric measures the effectiveness of our recommendations. A 5% rate is a conservative industry standard. This can be actively improved by partnering with trusted, reputable brands, ensuring a seamless user journey from diagnosis to purchase, and building user trust over time.
- **C (Average Commission):** This is the average revenue earned per sale. It will be determined through negotiations with brand partners and will depend on the price of the recommended products (e.g., pesticides, organic sprays, fertilizers). We assume a modest average of ₹40.

Multi-Year Forecast

By projecting growth in our user base and keeping other factors stable for a conservative estimate, we can forecast a growth trajectory.

Multi-Year Revenue Forecast for CropDoc

Metric	Year 1	Year 2	Year 3
Target Market Share	0.5%	1.2%	2.5%
Active Users (U)	750,000	1,800,000	3,750,000
Usage Frequency (F)	3	3	3
Conversion Rate (CR)	5%	5.5%	6%
Average Commission (C)	40 INR	42 INR	45 INR
Projected Revenue	₹45 Lakhs	₹1.2 Crores	₹2.8 Crores

Sensitivity Analysis

This financial model is highly sensitive to two key variables: **user acquisition (U)** and **conversion rate (CR)**. A faster user acquisition rate or a small 1-2% improvement in the conversion rate would lead to a significant, exponential increase in projected revenue. This highlights that the primary focus of the business should be on aggressive marketing to grow the user base and on optimizing the user experience to build the trust necessary for conversion.

Financial Equation for Market Growth

The overall financial model is built on the principle that revenue will grow as the user base expands over time. The total revenue is the sum of earnings from affiliate marketing, premium subscriptions, and data licensing.

The complete equation can be expressed as:

$$R_{total}(t) = U_0(1+g)^t \times [(F \times CR \times C) + (SR \times S_{price}) + D]$$

Component Breakdown (The Variables)

Here is a detailed explanation of each component in the equation:

- **Time-Based Growth (The Growth Engine):**
 - **$R_{total}(t)$** : The **Total Projected Revenue** at a specific time **t** .
 - **U_0** : The **Initial Number of Users** at the beginning ($t=0$).
 - **g** : The **Growth Rate** of the user base per time period (e.g., 50% annual growth would be **$g=0.5$**). This variable directly represents the market trend of increasing user adoption.
 - **t** : The **Time Period** (e.g., year 1, year 2, year 3).
- **Revenue Stream 1: Affiliate Marketing**
 - **F** : The average **Frequency** of diagnoses per user in that period.
 - **CR** : The **Conversion Rate** at which users purchase a sponsored product.
 - **C** : The average **Commission** earned per sale.
- **Revenue Stream 2: Premium Subscriptions**
 - **SR** : The **Subscription Rate** – the percentage of total users who convert to a paid subscription.
 - **S_{price}** : The **Average Annual Price** of a premium subscription.
- **Revenue Stream 3: Data Licensing**
 - **D** : The average **Data Revenue** generated per user per year from licensing anonymized data.

How the Equation Models the Trend

This equation is designed to correspond directly to the market trend of a growing user base.

The core of the model is the expression **$U_0(1+g)^t$** . This component calculates the compounding growth of your active users over several years. By multiplying this growing user base by the combined revenue generated from each user (from affiliates, subscriptions, and data), the model projects how your revenue will scale in alignment with the positive market trend.

This formula can be used in a spreadsheet to create detailed multi-year financial forecasts, allowing you to test different assumptions for growth rates, conversion rates, and pricing to see their impact on long-term profitability.




Figure 6

*** End of Report ***

Project Repository: <https://github.com/Chandra4243/CropDoc>

© 2025 Chandrabhushan Upadhyay. All Rights Reserved